Inside the loop

AI agents · Level 3 · Lesson 1



Today we will

- Read an agent’s conversation as one text with markers.
- Say who writes each part: the model or the app.
- Explain why every loop needs a stop.

Remember

- It’s all one text: tools, request, plan, call, result, answer.
- The model writes plans, calls and answers; the app adds the results.
- The loop stops at an answer, or at a step limit.

One long text

1. Warm-up (5 min): Where does an agent keep its plan?
2. Teach (10 min): The markers, and who writes what.
3. One long text (15 min): A conversation on a strip, in two colours.
4. Check it (5 min): Which colour can make mistakes?
5. Wrap-up (5 min): Where does the loop stop? Then the exit ticket.

Prepare a long paper strip (or A4 sheets taped together), blue and green pens or chalk, and the class timetable on the board. No phones are needed.

Let's discuss

1. Why is it safer that only the app runs tools and writes results?
2. What could go wrong in a loop with no step limit?
3. Where in the loop would you put a check for a message to parents?

Exit ticket

- Write the six parts of an agent’s text in order, and circle the ones the model writes.