Guardrails in code

AI agents · Upper secondary · Lesson 4

In this lesson you will

- Allow only the tools a task needs.
- Run an action only after a yes to that exact action.
- Check the details, and log every action.

Key ideas

- Learned safety can fail; the guardrail lives in code.
- An allow-list, exact approval, detail checks and a log.
- The model proposes; your code and a person decide.

Activity: Sneaky requests

1. Warm-up (3 min): Why did «نعم، أرسلها.» on its own get through?
2. Rules (6 min): Write four guardrails for a reminder agent.
3. Attack (6 min): A friend writes five sneaky requests.
4. Defend (5 min): Run, ask or block, for each one.

Sneaky requests stay harmless: they test the rules of a pretend agent, never a real account or a real person.

Let’s discuss

1. Why isn’t the model’s own “I’ll ask first” enough?
2. What should the log record, and who should read it?
3. Which tool would you never give a homework agent?

Exit ticket

- Write the four guardrail rules in your own words.