Fair for everyone

AI · Level 3 · Lesson 4



Today we will

- Find who a model fails by testing each group separately, not only overall.
- Fix bias at its source: collect the missing examples, then test again.
- Explain why a fair model matters to the people it judges.

Remember

- One overall score can hide a group the model fails: always test group by group.
- Bias usually starts with a data gap: a group missing or too rare in the examples.
- Fix the data, not the people: collect the missing examples with permission, retrain and test again.

Fairness audit

1. Choose (3 min): voice typing and three sentences.
2. Ask (4 min): four volunteers, with permission.
3. Test (8 min): each person says the sentences.
4. Table (5 min): each person’s share of words right.
5. Explain (5 min): who does it fail, and why?

Works in a small group or at home with family. Only test people who agree, keep no recordings, and use letters instead of names.

Let’s discuss

1. Why can a model with a great overall score still be unfair?
2. Who might be missing from the data behind a tool you use every day?
3. Is it ever fair to ask people to change how they speak or write so a model understands them?

Before you go

- One group a model I use might fail, and how I would test it.