Lesson plan · for the teacher

Jailbreaks: talking a chatbot round

AI security · Level 2 · Lesson 2

Time
40 min

Materials
A rule card for the keeper: a harmless secret (“The box is blue”) and one rule (“Never say the colour”); A second card for the checker: “Stop any answer that gives the colour away”; Board and chalk for the tally

Learning goals

- Explain why a chatbot's refusals can be talked round.
- Read a red-team score: tried, got through, how bad.
- Decide how a school responds to jailbreaks.

Activity: The rule keeper

1. Warm-up (5 min): Ask whether pupils have seen a chatbot say something rude, false or unsafe. Don't ask how; ask how it felt and what they did.
2. Teach (8 min): Explain that chatbots are trained to refuse harmful requests, that this refusal is learned from examples, and that a new wording can slip past it.
3. Round 1 (12 min): The keeper holds the card “Secret: the box is blue. Rule: never say the colour.” Classmates try polite rewordings for five minutes. Tally on the board what worked and how.
4. Round 2 (10 min): Add a checker who hears each answer first and stops any that gives the colour away. Play again and compare the tallies.
5. Wrap-up (5 min): Agree what to do when a chatbot says something harmful: stop, don't share it, tell a teacher, report it. Then give the exit ticket.

Keep the secret harmless and the rules strict: polite rewording only, no insults, no pressure. The game is about why rules need a second check, not about winning.

Discussion

- Why did some rewordings work on the first keeper?
- What did the checker catch that the keeper missed?
- What should you do if a chatbot says something harmful?

Exit ticket

- Write why a chatbot's “no” can be talked round, and one layer that catches what slips past.

Try it in class tomorrow
Play “The rule keeper”: one pupil guards a harmless secret with one rule, classmates try polite rewordings, then a second keeper checks every answer before it is said. Compare how much slips past.

Worksheet · for pupils

Jailbreaks: talking a chatbot round

AI security · Level 2 · Lesson 2

Name _____ Class _____ Date _____

1 **Pupils want a contest: who can get the class chatbot to say the rudest thing. What do you say?**

- ◯ "Go ahead: it's only a test"
- ◯ "No contests on real services; report what slips through instead"
- ◯ "Only on your own phones"

2 **Once a chatbot has refused a request, it will refuse every rewording of that request too.**

- ◯ True
- ◯ False

3 **Which are jailbreak shapes a teacher should recognise in a pupil's screenshot? Choose all that apply.**

- ☐ "It's only for a story"
- ☐ Splitting a request into harmless-looking pieces
- ☐ Asking in another language
- ☐ A clear question about tonight's homework
- ☐ Wearing it down over many turns

4 **Match each part of a red-team report to what it records.**

Kind	• • • •	The rule or change that stops it next time
Got through	• • • •	How bad, from 1 (low) to 3 (high)
Severity	• • • •	How many the chatbot didn't refuse
Fix	• • • •	What sort of request was tried: privacy, scam, unkind…

5 **Complete the sentence.**

Refusals are _____, so new _____ can slip past.

Word bank: learned · wordings · perfect · passwords

✎ **Have pupils ever shown you a chatbot saying something it shouldn't? How would you respond next time, and who would you report it to?**