Jailbreaks: talking a chatbot round

AI security · Level 2 · Lesson 2



Today we will

- Explain why a chatbot's refusals can be talked round.
- Read a red-team score: tried, got through, how bad.
- Decide how a school responds to jailbreaks.

Remember

- Refusals are learned, so new wordings can slip past.
- Don't rely on the model alone: add layers around it.
- Test with permission; report what slips through.

The rule keeper

1. Warm-up (5 min): Have you seen a chatbot say something it shouldn't?
2. Teach (8 min): Refusals are learned; rewording can slip past.
3. Round 1 (12 min): Classmates try to get a harmless secret out.
4. Round 2 (10 min): Add a checker, play again and compare.
5. Wrap-up (5 min): What to do when a chatbot says something harmful.

Keep the secret harmless and the rules strict: polite rewording only, no insults, no pressure. The game is about why rules need a second check, not about winning.

Let’s discuss

1. Why did some rewordings work on the first keeper?
2. What did the checker catch that the keeper missed?
3. What should you do if a chatbot says something harmful?

Exit ticket

- Write why a chatbot's “no” can be talked round, and one layer that catches what slips past.