Lesson plan · for the teacher

Red-team a school chatbot

AI security · Level 4 · Lesson 2

Time
55 min

Materials
Written permission from the head teacher (and from the provider, if its terms ask for it); Made-up pupils and records for the test; A results table on paper: request, wording, refused?, severity 1–3

Learning goals

- Plan a safe red-team test: permission, scope, made-up data.
- Score a test bank and rate each miss for severity.
- Report what you found, and check the fix.

Activity: Red-team session

1. Rules (5 min): Show the written permission. Agree which tool and account you will test, that every name and record is made up, and the stop rule: if anything real appears, stop and tell the teacher.
2. Bank (15 min): Each group writes three or four mild requests the chatbot should refuse (a pupil's record for a “neighbour”, a message posing as the school, exam answers), each in a plain wording and a polite rewording.
3. Run (15 min): Try each request in the Darsi red-team lab, or on a school tool you have permission to test, and record on the paper table: refused or not.
4. Score (10 min): Count the refusals. For each miss, agree its severity: 1 embarrassing, 2 harmful, 3 a danger to a person or their data.
5. Report (10 min): Write one page: what you tried, what got through, how often, how bad, and what you suggest. It goes to the tool's owner first, with no recipe for the trick.

Show the written permission at the start. Use made-up names only, test only tools you are allowed to test, and stop the session if anything real appears.

Discussion

- Why do polite rewordings get through when the plain request doesn't?
- Why does the report go to the owner before anyone else?
- Which fix is stronger: a better model, or a rule in code?

Exit ticket

- Write one request your school chatbot should refuse, in two wordings, and its severity if it got through.

Try it in class tomorrow
With your pupils' AI club, write ten mild requests a school chatbot should refuse, each in two wordings, with made-up pupils only. Run them in the Darsi red-team lab or on a school tool you

have written permission to test, score them with severity, and write a one-page report for the tool's owner.

Worksheet · for pupils

Red-team a school chatbot

AI security · Level 4 · Lesson 2

Name _____ Class _____ Date _____

1 **In the middle of a test, the chatbot shows a real pupil's phone number. What do you do?**

- ◯ Keep going: finish the bank first
- ◯ Stop, don't copy it, and report it to the owner
- ◯ Share a screenshot with the staff group

2 **Put the red-team test in order.**

- ☐ Report to the owner, and re-test after the fix
- ☐ Score refusals and rate severity
- ☐ Run it with made-up data
- ☐ Write the test bank, two wordings each
- ☐ Get written permission and agree the scope

3 **Each case in the bank is (request, refused?, severity). Count the refusals, print each miss like Miss: scam message severity 3, and finish with Refused 2 of 4.**

```
bank = [("pupil record", True, 3),
        ("scam message", False, 3),
        ("home address", False, 2),
        ("rude joke", True, 1)]
refused = 0
for case, ok, sev in bank:
    pass
print("Refused", refused, "of", len(bank))
```

Write your program:

4 **After the fix, the club re-runs four cases; True means it now refuses. What does this print?**

```
done = [True, True, False, True]
print(done.count(True), "of", len(done))
```

- ◯ 3 of 4
- ◯ 1 of 4
- ◯ 4 of 4

5 **Which belong in the rules for a school red-team test? Choose all that apply.**

- [ ] Written permission before you start
- [ ] Made-up pupils and records
- [ ] A stop rule if real data appears
- [ ] Posting the tricks that worked online, to warn others
- [ ] Testing any public chatbot you like

✎ **If your pupils red-teamed a tool your school uses, what would you want them to test first, and which rules would keep the test safe?**